


# Learning to learn by gradient descent by gradient descent



Presentation by Quan Bach

Paper Authors: Marcin Andrychowicz, Misha Denil, Sergio Gomez Colmenarejo, Matthew W. Hoffman, David Pfau, Tom Schaul, Brendan Shillingford, Nando de Freitas

# Introduction


## gradient descent

- Tasks in ML can be expressed as the problem of optimizing an objective function  $f(\theta)$  defined over some domain  $\theta \in \Theta$
- The goal is to find the minimizer  $\theta^* = \operatorname{argmin}_{\theta \in \Theta} f(\theta)$
- This can be solved by gradient descent if the function is differentiable, resulting in a sequence of updates:

$$\theta_{t+1} = \theta_t - \alpha \nabla f(\theta_t)$$

$\alpha$  : learning rate

$\nabla$  : gradient



# Introduction

## problem statement

- Optimization algorithms are still designed by hand
- Modern work in optimization is fragmented; designing update rules tailored to specific classes of problems
  - Deep learning community: high-dimensional, non-convex optimization problems (momentum, Rprop, Adagrad, RMSprop)
  - Other communities favor other approaches
- *No Free Lunch Theorems for Optimization*

specialization to a subclass of problems is in fact the only way that improved performance can be achieved in general

# Proposal

Replace hand-designed update rules with a learned update rule




Figure 1: The optimizer (left) is provided with performance of the optimizee (right) and proposes updates to increase the optimizee's performance. [photos: [Bobolas, 2009](#), [Maley, 2011](#)]

# Key Main Ideas

- Casting **algorithm design** as a learning problem
- A **learned update rule**

$$\theta_{t+1} = \theta_t + g_t(\nabla f(\theta_t), \phi), \quad g : \text{optimizer}, f : \text{optimizee}$$

- Model the update rules  $g$  using a **recurrent neural network (RNN)** which maintains its own state and hence dynamically updates as a function of its iterates

# Main Contribution of the Paper

## Methods

- Allowing for the output of back-propagation from one network to feed into an additional *learning* network, with both network trained jointly.
- Modification to the network architecture of the optimizer to scale to larger neural-network optimization problem
- Learning to learn a.k.a meta-learning with recurrent neural networks



# Main Contribution of the Paper

## Algorithms & Models

Take the update step  $g_t$  to be the output of a RNN  $m$ , parameterized by  $\phi$ , whose state will be denoted as  $h_t$

For some horizon time  $T$ , given a distribution of functions  $f$  the expected loss is written as:

$$\mathcal{L}(\phi) = E_f \left[ \sum_{t=1}^T w_t f(\theta_t) \right]; \text{ where } \theta_{t+1} = \theta_t + g_t$$
$$\begin{bmatrix} g_t \\ h_{t+1} \end{bmatrix} = m(\nabla_t, h_t, \phi)$$

$w_t \in \mathbb{R}_{\geq 0}$  : are arbitrary weights associated with each time step

Note :  $\nabla_t = \nabla_{\theta} f(\theta_t)$

Minimizing the value of  $\mathcal{L}(\phi)$  using gradient descent on  $\phi$

# Main Contribution of the Paper

## Computational Graph

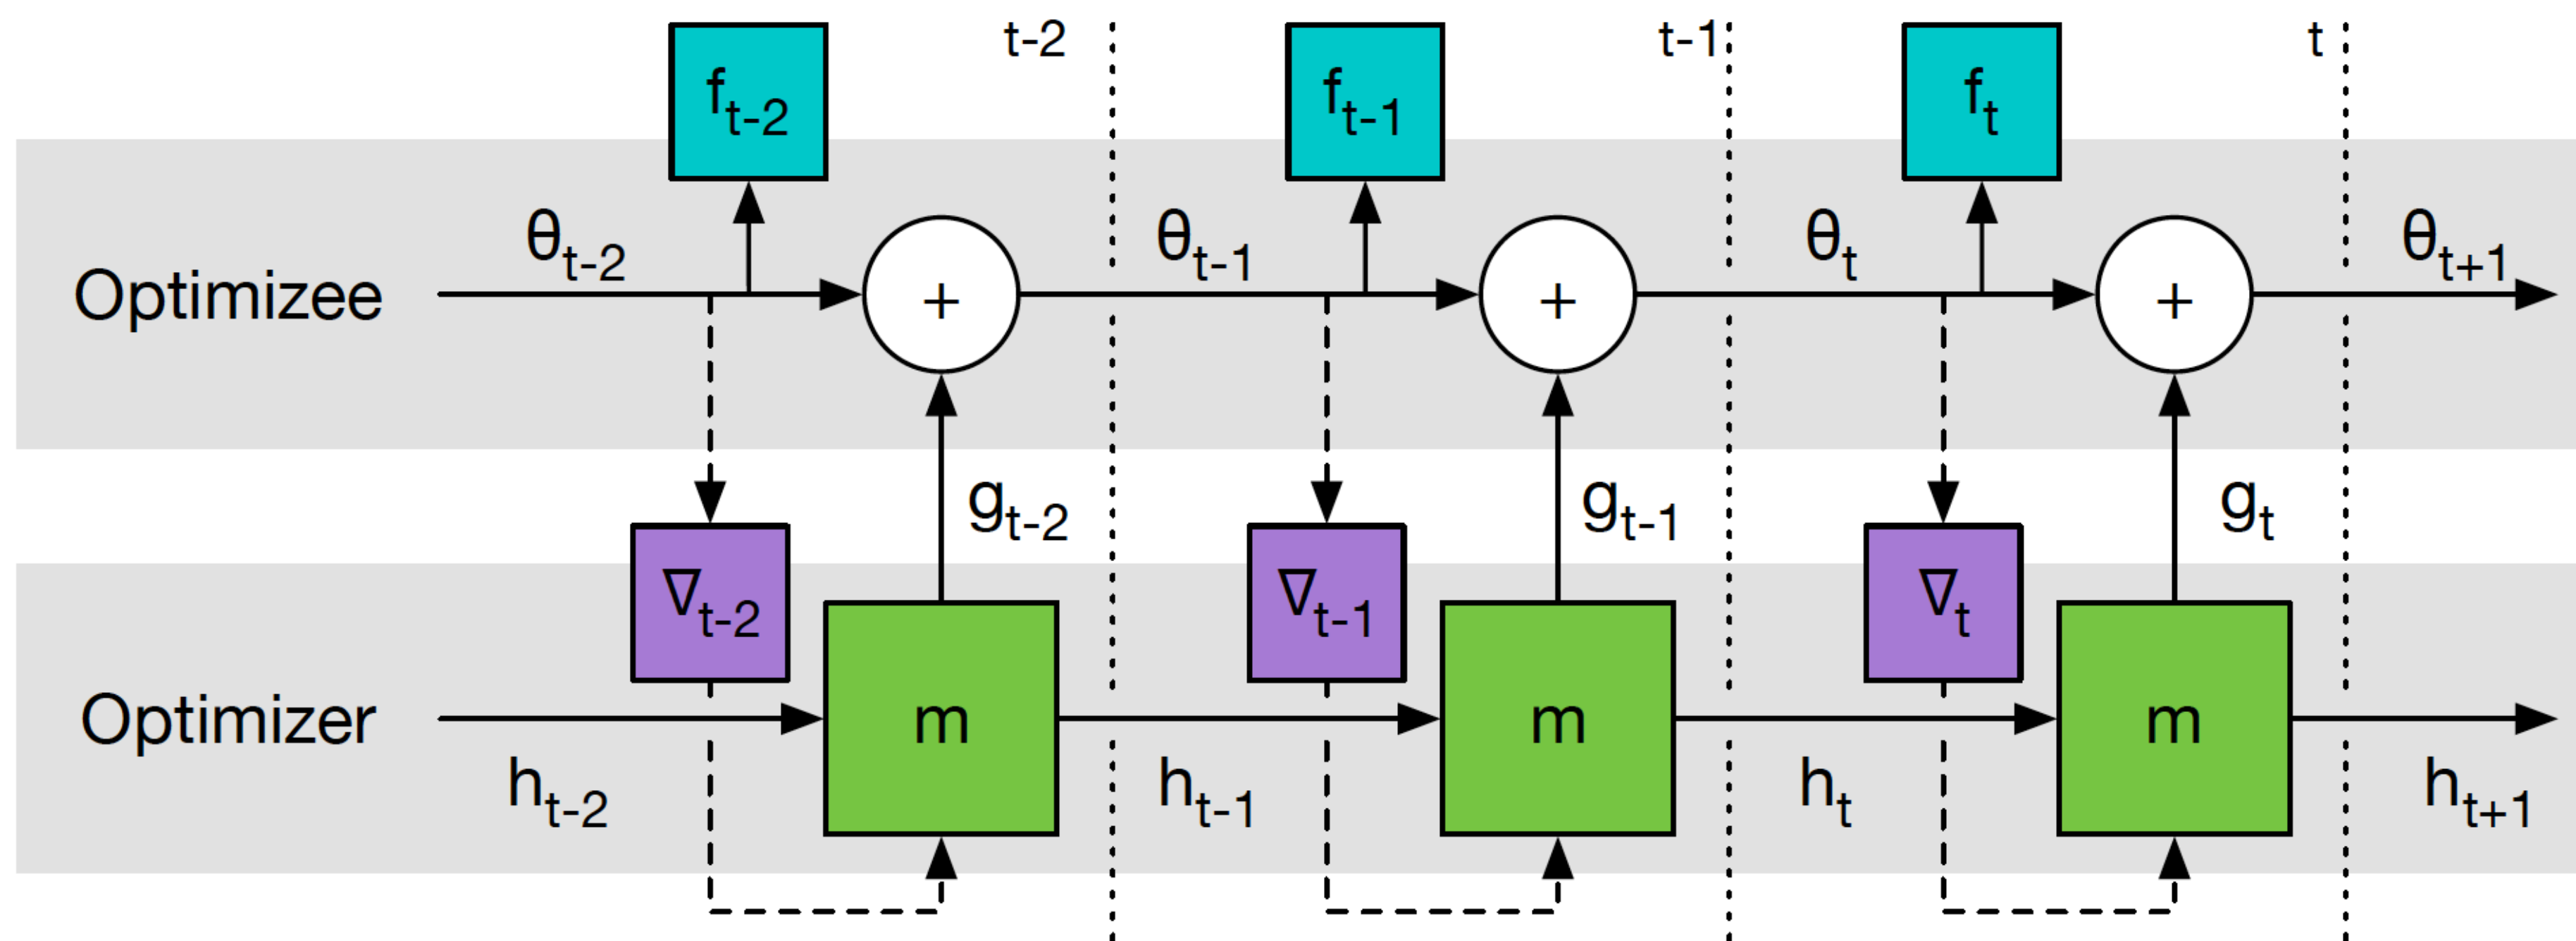


Figure 2: Computational graph used for computing the gradient of the optimizer.



# Main Contribution of the Paper

## Long Short Term Memory (LSTM) optimizer

- RNNs have at least tens of thousands of parameters
- Solution: optimizer  $m$  operates coordinate-wise on the parameters  
(*coordinatewise network architecture*)
- LSTM: is a RNN which has feedback connections

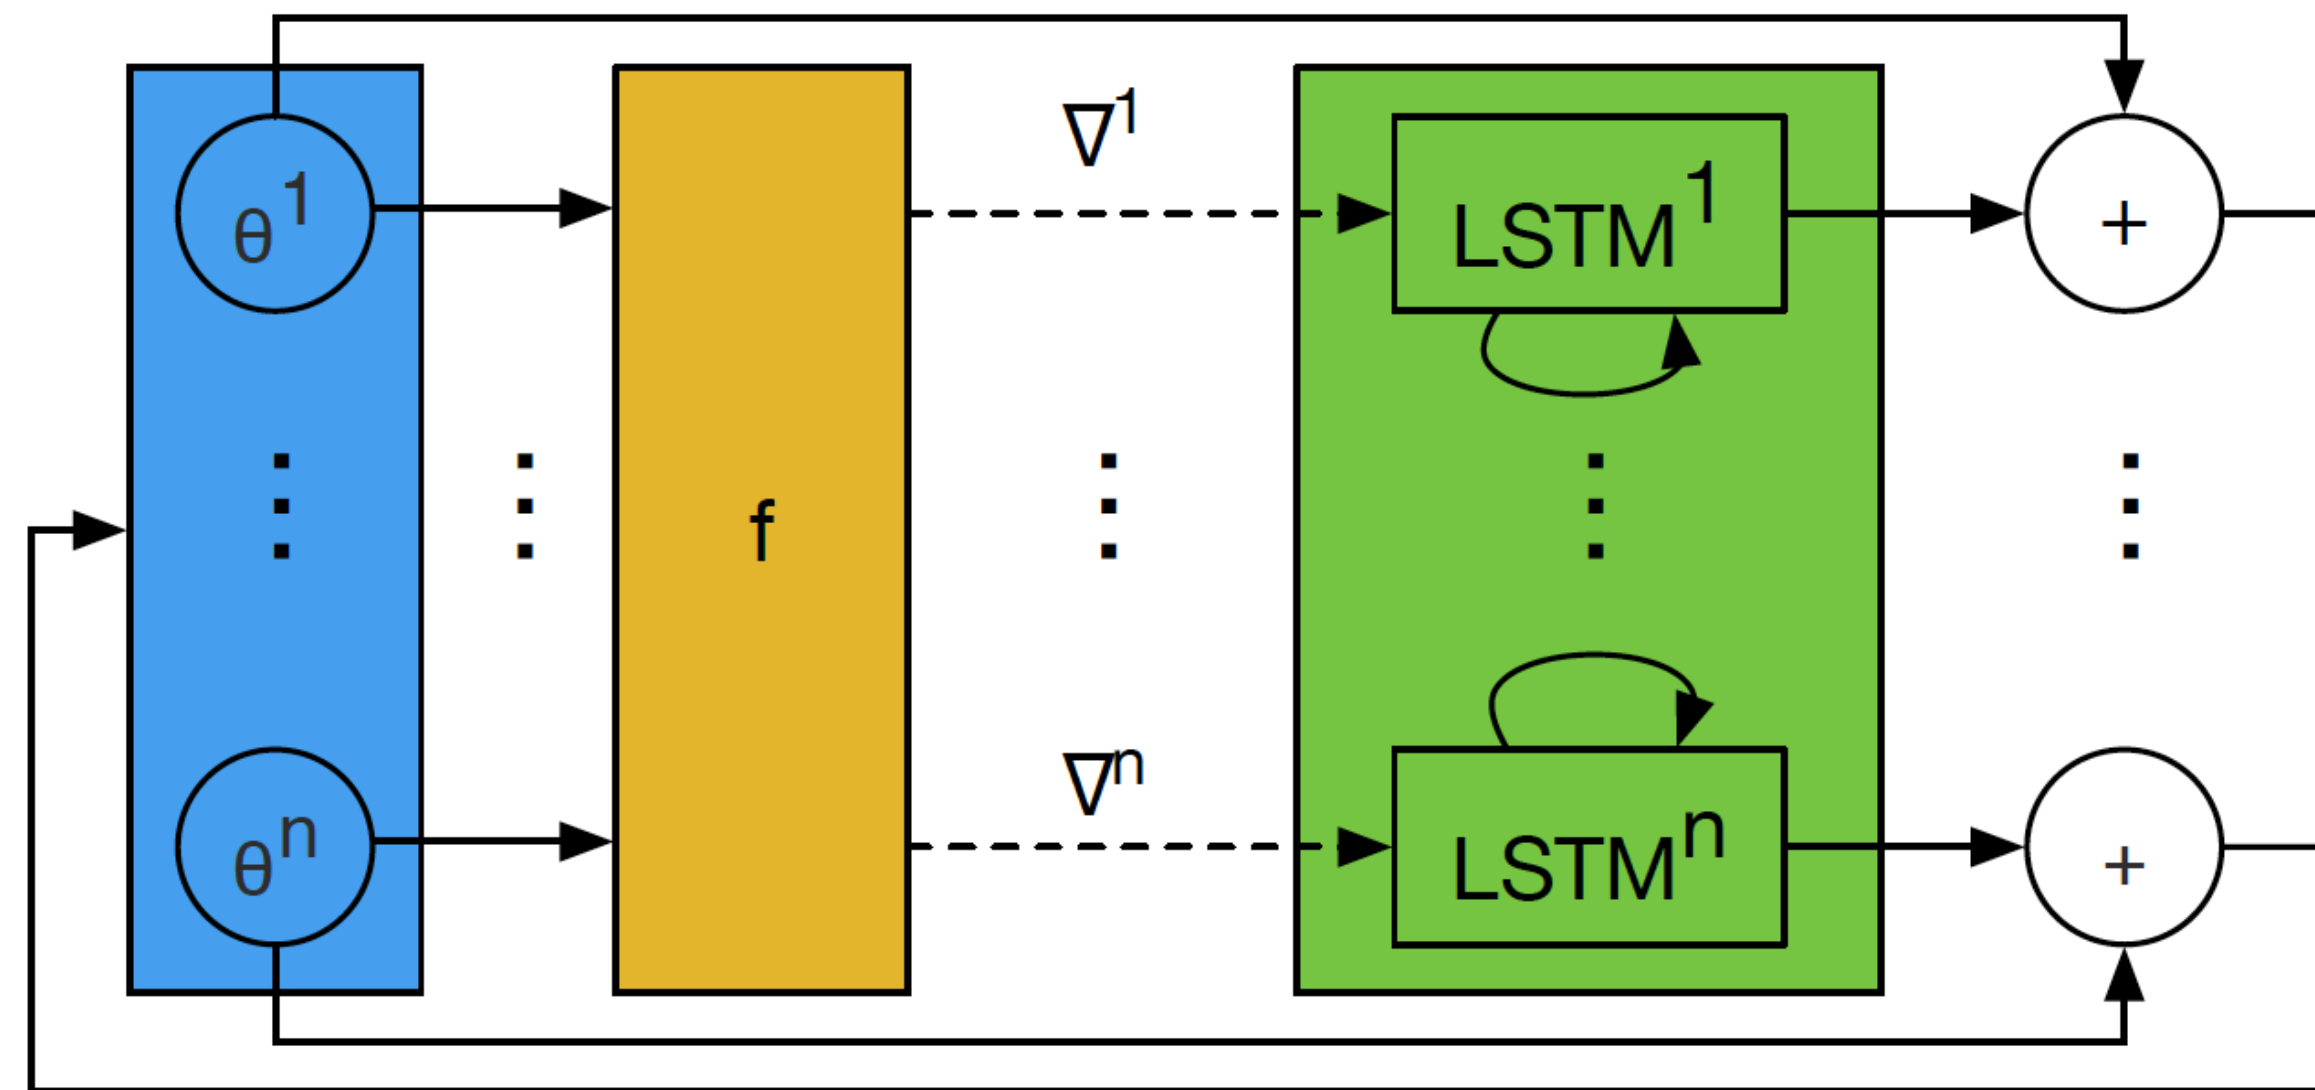


Figure 3: One step of an LSTM optimizer. All LSTMs have shared parameters, but separate hidden states.

# Evaluation Design

Comparison with standard optimizers:

- Stochastic Gradient Descent
- RMSprop
- ADAM
- Nesterov's accelerated gradient (NAG)

Datasets:

- Quadratic functions:  $f(\theta) = ||W\theta - y||_2^2$ ; with 10x10 matrices W and 10-dimensional vector y
- MNIST
- CIFAR-10
- ImageNet

# Experimental Results

## Quadratic functions & MNIST




Figure 4: Comparisons between learned and hand-crafted optimizers performance. Learned optimizers are shown with solid lines and hand-crafted optimizers are shown with dashed lines. Units for the  $y$  axis in the MNIST plots are logits. **Left:** Performance of different optimizers on randomly sampled 10-dimensional quadratic functions. **Center:** the LSTM optimizer outperforms standard methods training the base network on MNIST. **Right:** Learning curves for steps 100-200 by an optimizer trained to optimize for 100 steps (continuation of center plot).



# Experimental Results

## MNIST




Figure 5: Comparisons between learned and hand-crafted optimizers performance. Units for the  $y$  axis are logits. **Left:** Generalization to the different number of hidden units (40 instead of 20). **Center:** Generalization to the different number of hidden layers (2 instead of 1). This optimization problem is very hard, because the hidden layers are very narrow. **Right:** Training curves for an MLP with 20 hidden units using ReLU activations. The LSTM optimizer was trained on an MLP with sigmoid activations.

# Experimental Results

## CIFAR




Figure 7: Optimization performance on the CIFAR-10 dataset and subsets. Shown on the left is the LSTM optimizer versus various baselines trained on CIFAR-10 and tested on a held-out test set. The two plots on the right are the performance of these optimizers on subsets of the CIFAR labels. The additional optimizer *LSTM-sub* has been trained only on the heldout labels and is hence transferring to a completely novel dataset.



# Experimental Results

## Neural Art

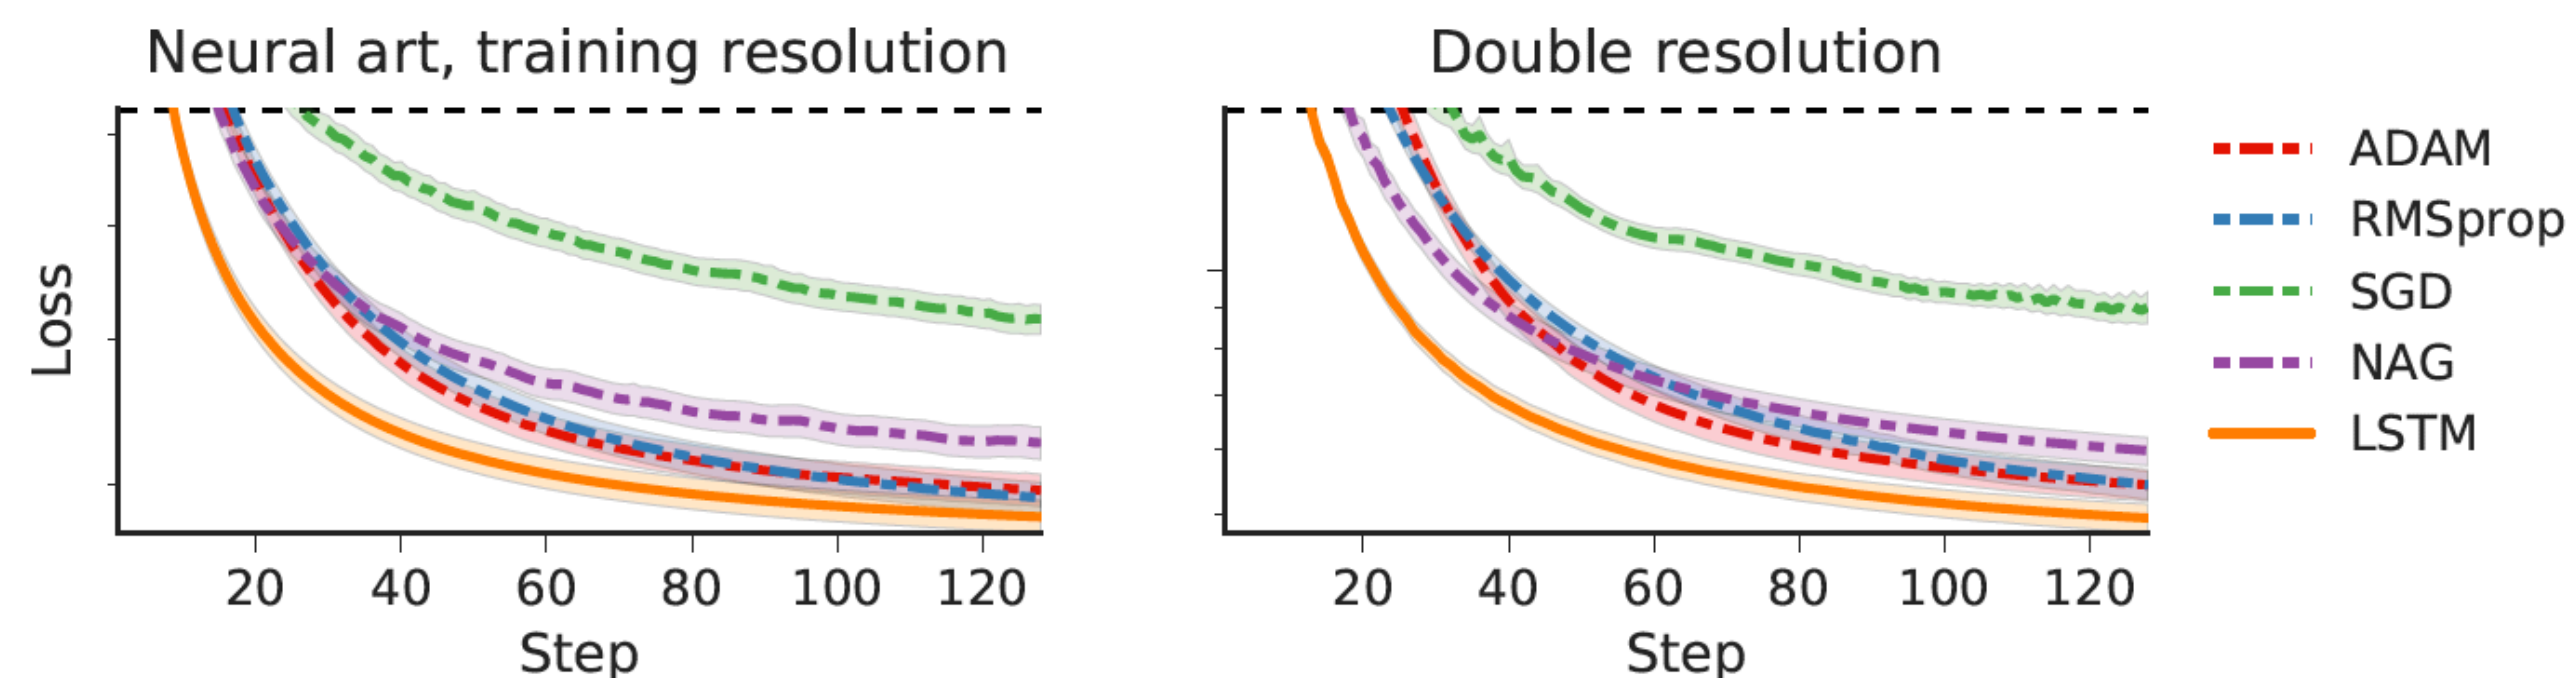


Figure 8: Optimization curves for Neural Art. Content images come from the test set, which was not used during the LSTM optimizer training. Note: the y-axis is in log scale and we zoom in on the interesting portion of this plot. **Left:** Applying the training style at the training resolution. **Right:** Applying the test style at double the training resolution.




Figure 9: Examples of images styled using the LSTM optimizer. Each triple consists of the content image (left), style (right) and image generated by the LSTM optimizer (center). **Left:** The result of applying the training style at the training resolution to a test image. **Right:** The result of applying a new style to a test image at double the resolution on which the optimizer was trained.



# Critique the main contribution

## Challenges with RNNs

- Computation being slow
- Difficulty accessing information from a long time ago

### Critique:

- **No** comparison in the time it takes to train the network
- **No** comparison in the time to run the network (only how the loss function change each step)

# Critique the main contribution

## Method of training

Early stopping is implemented while training; the authors would stop after each epoch, freeze the parameters and evaluate its performance.

### Critique:

- This method of training is very **time consuming**
- Required **human evaluation**

# Critique the main contribution

## Models to compare

The authors tuned the learning rate for SGD, RMSProp, ADAM, and NAG

Critique:

- These standard optimizers **might not be at their optimal performance**

# Discuss Questions

## Update Rule

Standard optimizer's sequence update:

$$\theta_{t+1} = \theta_t - \alpha \nabla f(\theta_t) \quad (1)$$

LSMT optimizer's sequence update:

$$\theta_{t+1} = \theta_t + g_t(\nabla f(\theta_t), \phi) \quad (2)$$

What is the significant of the (-) in equation (1)?

Why does it change to the (+) in (2)?